

Assignments Covered

- Working with AWS Databases and Lambda: Postgres (Base Task)
- Working with AWS Databases and Lambda: Postgres (Bonus 3)
- Working with AWS Databases and the Command Line/Terminal: DynamoDB (Base Task)

(The tutorial was also finished, but that was simply following the instructions provided)

All work was done in region us-west-2(Oregon) whenever possible.

####

Working with AWS Databases and Lambda: Postgres (Base Task)

I created the RDS database by following the tutorial provided and called it “rds-postgres-tutorial” with the master username “postgres” and the password “postgres.”

Afterwards, I set up two lambda functions: one to stop the server, and one to start it. The process for creating both is somewhat similar. I will go over the process for setting up the lambda function to stop the server and cover the differences between it and the function to start the server.

I created an IAM policy called “RDS_Start_Stop_Policy” that gave permissions to check for existing databases on my account as well as permissions to start/stop database instances using the following JSON:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "rds:DescribeDBInstances",
        "rds:StopDBInstance",
        "rds:StartDBInstance"
      ],
      "Resource": "*"
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:CreateLogGroup",
        "logs:CreateLogStream",
        "logs:PutLogEvents"
      ],

```

```

        "Resource": "*"
    }
]
}

```

After that, I created the role “AWSLambdaRole” in the IAM Management Console. This role can be used with Lambda functions and utilizes “RDS_Start_Stop_Policy.”

==(From here, the function to start the database undergoes a similar, but separate process)==

Next, I created the Lambda function “RDSInstanceStop.” This function utilizes Python 3.8 and uses the IAM Role “AWSLambdaRole” for its permissions. Using the ARN of the function after creation, I also created a new inline policy called “policy_rds_stop” using the following JSON code:

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "lambda:GetFunctionConfiguration",
      "Resource":
"arn:aws:lambda:us-west-2:367829174649:function:RDSInstanceStop"
    }
  ]
}

```

The string following “Resource” is the ARN of “RDSInstanceStop.”

After that, I turned towards the code of the Lambda function itself. I used the following Python 3.8 code:

```

import sys
import boto3
import boto3
from boto3.exceptions import ClientError

def lambda_handler(event, context):
    rds = boto3.client('rds')
    lambdaFunc = boto3.client('lambda')
    print ('Trying to get Environment variable')
    try:
        funcResponse = lambdaFunc.get_function_configuration(
            FunctionName='RDSInstanceStop'
        )
        DBInstance = funcResponse['Environment']['Variables']['DBInstanceName']
        print ('Stopping RDS service for DBInstance : ')
    except ClientError as e:
        print(e)
    try:
        response = rds.stop_db_instance(

```

```

        DBInstanceIdentifier=DBInstance
    )
    print ('Success')
except ClientError as e:
    print(e)

    return {
        'message' : "Script execution completed. See Cloudwatch logs for complete
output"
    }

```

After deploying this code, I created a default test and ran said test. The result is that the server enters a “Stopping” status for a few minutes before finally stopping.

To create the Lambda function that starts the server back up again, I used a similar process to create the function “RDSInstanceStart.” Slight changes of the JSON and Python 3.8 code can be found in the zipped folder.

###

Working with AWS Databases and Lambda: Postgres (Bonus 3)

For this task, I recreated an instance of the database “rds-postgres-tutorial” in CloudFormation by using the CloudFormation console.

First, I selected to create a new stack using existing resources. I was prompted to provide a template for this stack, which is shown below:

```

{
  "Resources" : {
    "MyDB" : {
      "Type" : "AWS::RDS::DBInstance",
      "Properties" : {
        "AllocatedStorage" : "20",
        "DBInstanceClass" : "db.t2.micro",
        "Engine" : "PostgreSQL",
        "MasterUsername" : "postgres",
        "MasterUserPassword" : "postgres"
      },
      "DeletionPolicy" : "Snapshot"
    }
  }
}

```

The information in “Properties” was gathered from the overview of “rds-postgres-tutorial” in the RDS console.

The other major thing I did for setup is referring the stack to “rds-postgres-tutorial” for duplication. Upon finishing the setup, the stack began importing and updating the information.

Looking back, one way I could have improved this replication was to include the security group used in “rds-postgres-tutorial.” Since I didn’t declare what security group I was using, it replaced the new security group I set up using the tutorial to a default security group. This specific update took a somewhat long time (about 20 minutes) and I believe I could have avoided this if I specified the security group in the template.

###

Working with AWS Databases and the Command Line/Terminal: DynamoDB (Base Task)

After installing AWS CLI, I went on to the IAM Management Console and added a user “ayhu0311” with admin privileges, Access Key, and Secret Access Key. From there, I configured that user using **aws configure** in the terminal.

Once I’ve finished setting up the user “ayhu0311” as the base user on the command line, I created a new table in DynamoDB using the following command line input:

```
aws dynamodb create-table \  
--table-name Shops \  
--attribute-definitions \  
AttributeName=ShopName,AttributeType=S \  
--key-schema \  
AttributeName=ShopName,KeyType=HASH \  
--provisioned-throughput \  
ReadCapacityUnits=10,WriteCapacityUnits=5
```

This creates the table “Shops” with the primary key attribute “ShopName” that utilizes a hash index.

After creating the table, I inserted an item using the following command line input:

```
aws dynamodb put-item \  
--table-name Shops \  
--item '{  
"ShopName": {"S": "Shop1"},  
"Time": {"N": "15"},  
"Price": {"N": "15.00"}  
}'
```

This inserts a row with the attributes “ShopName,” “Time,” and “Price.”

(The practical use for this table would be to insert a bunch of different shops that all sell some one particular item/service of interest at various prices and have different travel times to reach.)

Once I've inserted an item, I retrieved it using the following command line input:

```
aws dynamodb get-item \  
--table-name Shops --key '{"ShopName": {"S":"Shop1"}}'
```

###

References

Besides official documentation, I also referred to the following link to set up the lambda functions for starting/stopping the database:

<https://www.sqlshack.com/automatically-start-stop-an-aws-rds-sql-server-using-aws-lambda-functions/>